



Copyright © 2006 Intel Corporation. All Rights Reserved.

Using the Pin Instrumentation Tool for Computer Architecture Research

Aamer Jaleel, Chi-Keung Luk, Bobbie Manne, Harish Patil

Intel® Corporation

June 17, 2006

Introduction to Pin

- 1 -

Copyright © 2006 Intel Corporation. All Rights Reserved.

What is Instrumentation?

- A technique that inserts extra code into a program to collect runtime information
- Instrumentation approaches:
 - Source instrumentation:
 - Instrument source programs
 - **Binary instrumentation:** ← our approach
 - Instrument executables directly

- 2 -

Example: Instruction Count

```
counter++;  
sub $0xff, %edx  
counter++;  
cmp %esi, %edx  
counter++;  
jle <L1>  
counter++;  
mov $0x1, %edi  
counter++;  
add $0x10, %eax
```

- 3 -

Example: Instruction Trace

```
Print(ip);  
sub $0xff, %edx  
Print(ip);  
cmp %esi, %edx  
Print(ip);  
jle <L1>  
Print(ip);  
mov $0x1, %edi  
Print(ip);  
add $0x10, %eax
```

- 4 -

Instrumentation vs. Simulation

- Advantages of Simulation:
 - Detailed modeling of processors
 - Can model non-existing hardware
- Advantages of Instrumentation:
 - Easy to prototype
 - Fast to run (allowing complete runs)

- 5 -

How is Instrumentation used in Computer Architecture?

- Trace Generation
- Branch Predictor and Cache Modeling
- Fault Tolerance Study
- Emulating Speculation
- Emulating New Instructions

- 6 -

What is Pin?

- **Easy-to-use Instrumentation:**
 - Uses dynamic instrumentation
 - Do not need source code, recompilation, post-linking
- **Programmable Instrumentation:**
 - Provides rich APIs to write in C/C++ your own instrumentation tools (called **Pintools**)
- **Multiplatform:**
 - Supports IA-32, EM64T, Itanium, Xscale
 - Supports Linux, Windows, MacOS
- **Robust:**
 - Instruments real-life applications
 - Database, search engines, web browsers, ...
 - Instruments multithreaded applications
- **Efficient:**
 - Applies compiler optimizations on instrumentation code

- 7 -

How to use Pin?

- Launch and instrument an application

```
$ pin -t pintool -- application
```

Instrumentation engine
(provided in our kit)

Instrumentation tool
(write your own, or use one
provided in our kit)

- Attach to and instrument an application

```
$ pin -t pintool -pid 1234
```

- 8 -

Writing Pintools

- 9 -

Pin Instrumentation APIs

- Basic APIs are architecture independent:
 - Provide common functionalities like determining:
 - Control-flow changes
 - Memory accesses
- Architecture-specific APIs
 - E.g., Info about segmentation registers on IA32
- Call-based APIs:
 - Instrumentation routines
 - Analysis routines

- 10 -

Instrumentation vs. Analysis

Concepts borrowed from the ATOM tool:

- **Instrumentation routines** define where instrumentation is **inserted**
 - e.g. before instruction
 - ☞ **Occurs *first time* an instruction is executed**
- **Analysis routines** define what to do when instrumentation is **activated**
 - e.g. increment counter
 - ☞ **Occurs *every time* an instruction is executed**

- 11 -

Pintool 1: Instruction Count

```
counter++;  
sub $0xff, %edx  
counter++;  
cmp %esi, %edx  
counter++;  
jle <L1>  
counter++;  
mov $0x1, %edi  
counter++;  
add $0x10, %eax
```

- 12 -

Pintool 1: Instruction Count Output

```
$ /bin/ls
Makefile atrace.o imageload.out itrace
proccount Makefile.example imageload
inscount0 itrace.o proccount.o atrace
imageload.o inscount0.o itrace.out
```

```
$ pin -t inscount0 -- /bin/ls
Makefile atrace.o imageload.out itrace
proccount Makefile.example imageload
inscount0 itrace.o proccount.o atrace
imageload.o inscount0.o itrace.out
```

Count 422838

- 13 -

ManualExamples/inscount0.C

```
#include <iostream>
#include "pin.h"
```

- Same source code works on the 4 architectures
- Pin automatically and efficiently saves/restores application state

```
{ std::cerr << "Count " << icount << endl; }

int main(int argc, char * argv[])
{
    PIN_Init(argc, argv);
    INS_AddInstrumentFunction(Instruction, 0);
    PIN_AddFiniFunction(Fini, 0);
    PIN_StartProgram();
    return 0;
}
```

- 14 -

Pintool 2: Instruction Trace

- Need to pass an argument (ip) to the analysis routine (printip())

```
Print(ip);
mov $0x1, %edi

Print(ip);
add $0x10, %eax
```

- 15 -

Pintool 2: Instruction Trace Output

```
$ pin -t itrace -- /bin/ls
Makefile atrace.o imageload.out itrace
proccount Makefile.example imageload
inscount0 itrace.o proccount.o atrace
imageload.o inscount0.o itrace.out
```

```
$ head -4 itrace.out
```

```
0x40001e90
```

```
0x40001e91
```

```
0x40001ee4
```

```
0x40001ee5
```

- 16 -

ManualExamples/itrace.C

```

#include <stdio.h>
#include "pin.H"
FILE * trace;
void printip(void *ip) { fprintf(trace, "%p\n", ip); }
void Instruction(INS ins, void *v) {
    INS_InsertCall(ins, IPOINT_BEFORE, (AFUNPTR)printip,
        IARG_INST_PTR, IARG_END);
}
void Fini(INT32 code, void *v) { fclose(trace); }
int main(int argc, char * argv[]) {
    trace = fopen("itrace.out", "w");
    PIN_Init(argc, argv);
    INS_AddInstrumentFunction(Instruction, 0);
    PIN_AddFiniFunction(Fini, 0);
    PIN_StartProgram();
    return 0;
}

```

argument to analysis routine

analysis routine

instrumentation routine

- 17 -

Examples of Arguments to Analysis Routine

- **IARG_INST_PTR**
 - Instruction pointer (program counter) value
- **IARG_UINT32 <value>**
 - An integer value
- **IARG_REG_VALUE <register name>**
 - Value of the register specified
- **IARG_BRANCH_TARGET_ADDR**
 - Target address of the branch instrumented
- **IARG_MEMORY_READ_EA**
 - Effective address of a memory read

And many more ... (refer to the Pin manual for details)

Instrumentation Points

■ Instrument points relative to an instruction:

- Before (*IPOINT_BEFORE*)
- After:
 - Fall-through edge (*IPOINT_AFTER*)
 - Taken edge (*IPOINT_TAKEN*)

```

count() → cmp    %esi, %edx
count() → jle    <L1>
count() → mov    $0x1, %edi
                                mov    $0x8, %edi
                                <L1>:

```

- 19 -

Instrumentation Granularity

■ Instrumentation with Pin can be done at 3 different granularities: 1 Trace, 2 basic blocks, 6 insts

- Instruction
- Basic block
 - A sequence of instructions (conditional or unconditional)
 - Single entry, single exit
- Trace
 - A sequence of basic blocks (conditional or unconditional control-flow changing instruction)
 - Single entry, multiple exits

```

sub    $0xff, %edx
cmp    %esi, %edx
jle    <L1>

mov    $0x1, %edi
add    $0x10, %eax
jmp    <L2>

```

- 20 -

Recap of Pintool 1: Instruction Count

counter++;

- Straightforward, but the counting can be more efficient

```
counter++;
jle <L1>

counter++;
mov $0x1, %edi

counter++;
add $0x10, %eax
```

- 21 -

Pintool 3: Faster Instruction Count

```
counter += 3
sub $0xff, %edx

cmp %esi, %edx

jle <L1>

counter += 2
mov $0x1, %edi

add $0x10, %eax
```

basic blocks (bb1)

- 22 -

```
#include <stdio.h>
```

```
#include "pin.H"
```

```
UINT64 icount = 0;
```

```
void docount(INT32 c) { icount += c; }
```

analysis routine

```
void Trace(TRACE trace, void *v) {
```

instrumentation routine

```
    for (BBL bbl = TRACE_EblHead(trace);
```

```
        BBL_Valid(bbl); bbl = BBL_Next(bbl)) {
```

```
        BBL_InsertCall(bbl, IPOINT_BEFORE, (AFUNPTR)docount,
```

```
                        IARG_UINT32, BBL_NumIns(bbl), IARG_END);
```

```
    }
```

```
}
```

```
void Fini(INT32 code, void *v) {
```

```
    fprintf(stderr, "Count %lld\n", icount);
```

```
}
```

```
int main(int argc, char * argv[]) {
```

```
    PIN_Init(argc, argv);
```

```
    TRACE_AddInstrumentFunction(Trace, 0);
```

```
    PIN_AddFiniFunction(Fini, 0);
```

```
    PIN_StartProgram();
```

```
    return 0;
```

```
}
```

- 23 -

Copyright © 2006 Intel Corporation. All Rights Reserved.

Modifying Program Behavior

- Pin allows you not only observing but also changing program behavior
- Ways to change program behavior:
 - Add/delete instructions
 - Change register values
 - Change memory values
 - Change control flow

- 24 -

Example: Emulation of Loads

```
sub    $0x11c,%esp  
  
mov    0xc(%ebp),%eax  
  
add    $0x128,%eax  
  
mov    0x8(%ebp),%edi  
  
xor    %eax,%edi
```

- 25 -

Example: Emulation of Loads

```
sub    $0x11c,%esp  
EmulateLoad(%ebp+0xc,%eax)  
mov    0xc(%ebp),%eax  
  
add    $0x128,%eax  
EmulateLoad(%ebp+0x8,%edi)  
mov    0x8(%ebp),%edi  
  
xor    %eax,%edi
```

- 26 -

Emulation of Loads

```
$ pin -t emuload -- /bin/ls
```

```
Emulate loading from addr 0xbffffe188
Emulate loading from addr 0x40016ae0
Emulate loading from addr 0x40016c74
Emulate loading from addr 0x40016c7c
Emulate loading from addr 0x40016c84
...
_insprofiler.C imageload      imageload.out
insprofiler.C  proccount.C atrace.C
```

- 27 -

SimpleExamples/emuload.C

```
#include <stdio.h>
#include "pin.H"
#include "pin_isa.H"
#include <iostream>
```

```
// Move from memory to register
ADDRINT DoLoad(ADDRINT * addr) {
    cout << "Emulate loading from addr " << addr << endl;
    return *addr;
}
```

```
VOID EmulateLoad(INS ins, VOID* v) {
    if (INS_Opcode(ins) == XEDICLASS_MOV && INS_IsMemoryRead(ins) &&
        INS_OperandIsReg(ins, 0) && INS_OperandIsMemory(ins, 1)) {
        // op0 <- *op1
        INS_InsertCall(ins, IPPOINT_BEFORE, AFUNPTR(DoLoad),
            IARG_MEMORYREAD_EA,
            IARG_RETURN_REGS, INS_Op0,
            IARG_END);
        INS_Delete(ins);
    }
}
```

check if ins is a load
pass the load data address to DoLoad()
use DoLoad()'s return value to remove ins (the original load) register which is the first operand of ins

```
int main(int argc, char * argv[]) {
    PIN_Init(argc, argv);
    INS_AddInstrumentFunction(EmulateLoad, 0);
    PIN_StartProgram();
    return 0;
}
```

- 28 -

Multithreading Support

- Notify the pintool when a thread is created or exited
- Provide a "thread id" for pintools to identify a thread
- Provide locks for pintools to access shared data structures

- 29 -

Example of Instrumenting Multithreaded Programs

```
$ pin -mt -t mtest -- thread
Creating thread
Creating thread
Joined 0
Joined 1
$ cat mtest.out
0x400109a8: 0
thread begin 1 sp 0x80acc00 flags f00
0x40001d38: 1
thread begin 3 sp 0x43305bd8 flags f21
0x40011220: 3
thread begin 2 sp 0x42302bd8 flags f21
0x40010e15: 2
0x40005cdc: 2
thread end 3 code 0
0x40005e90: 0
0x40005e90: 0
thread end 2 code 0
thread end 1 code 0
```

- 30 -

FILE * out;
PIN_LOCK lock;

Copyright © 2006 Intel Corporation. All Rights Reserved.
[Tests/mtest.C](#)

analysis routine

```

VOID TraceBegin(VOID * ip, UINT32 threadid) {
    GetLock(&lock, threadid+1);
    fprintf(out, "%p: %d\n", ip, threadid);
    ReleaseLock(&lock);
}

VOID Trace(TRACE trace, VOID *v) {
    TRACE_InsertCall(trace, IPOINT_BEFORE, AFUNPTR(TraceBegin),
        IARG_INST_PTR, IARG_THREAD_ID, IARG_END);
}

VOID ThreadBegin(UINT32 threadid, VOID * sp, int flags, VOID *v) {
    GetLock(&lock, threadid+1);
    fprintf(out, "thread begin %d sp %p flags %x\n", threadid, sp, flags);
    ReleaseLock(&lock);
}

VOID ThreadEnd(UINT32 threadid, INT32 code, VOID *v) {
    GetLock(&lock, threadid+1);
    fprintf(out, "thread end %d code %d\n", threadid, code);
    ReleaseLock(&lock);
}

VOID Fini(INT32 code, VOID *v) {
    fprintf(out, "Fini: code %d\n", code);
}

int main(INT32 argc, CHAR **argv) {
    InitLock(&lock);
    out = fopen("mtest.out", "w");
    PIN_Init(argc, argv);
    PIN_AddThreadBeginFunction(ThreadBegin, 0);
    PIN_AddThreadEndFunction(ThreadEnd, 0);
    TRACE_AddInstrumentFunction(Trace, 0);
    PIN_AddFiniFunction(Fini, 0);
    PIN_StartProgram();
    return 0;
}

```

instrumentation routines

- 31 -

Debugging Pintools

1. Invoke gdb with your pintool (but don't use "run")

```

$ gdb inscount0
(gdb)

```

2. On another window, start your pintool with "-pause_tool"

```

$ pin -pause_tool 5 -t inscount0 -- /bin/ls
Pausing to attach to pid 32017

```

3. Go back to gdb:

- a) Attach to the process
- b) Use "cont" to continue execution; can set breakpoints as usual

```

(gdb) attach 32017
(gdb) break main
(gdb) cont

```

- 32 -

Conclusions

Pin

- Build your own architectural tools with ease
- Run on multiple platforms:
 - IA-32, EM64T, Itanium, and XScale
 - Linux, Windows, MacOS
- Work on real-life applications
- Efficient instrumentation

- 33 -

Call For Action

Try it out!

Free download from:

<http://roque.colorado.edu/pin>

* User manual, many example tools, tutorials

* ~7000 downloads since 2004 July

Group: [**http://groups.yahoo.com/group/pinheads/**](http://groups.yahoo.com/group/pinheads/)

Email: [**pin.project@intel.com**](mailto:pin.project@intel.com)

- 34 -